Sentiment Analysis of Yelp Reviews Using Transformers

Kalra Srishti
Masters of Science Computer Science
Texas A&M University
College Station, Texas
srishtikalra@tamu.edu

***Abstract*—As a part of the final project for CSCE 633 Machine Learning Course, I have implemented a transformer model that uses Bidirectional Encoder Representations from Transformers (BERT) for sentiment analysis based on text review and stars. The BERT model has been modified to perform this task. A very important aspect of this project involves experimenting with different hyper parameters and custom layers to fine-tune this pre-trained model. This process is critical for the BERT to understand the specific linguistic characteristics and sentiment expressions found in the data. To achieve a seamless integration of the dataset with the model, I have developed specialized data loaders to process and align the review data with the BERT architecture. A key innovation in the methodology used is the modification of BERT's final layers, a strategic decision aimed at enhancing the model's ability to discern and classify sentiments as either positive, negative, or neutral. This fine-tuning process, along with the architectural adjustments, was very carefully implemented through a series of training loops. This helped in balancing the pre-trained model's broad language understanding with the specific task requirements. The efficacy of this approach was rigorously validated using standard evaluation metrics, where the modified BERT model demonstrated a marked improvement in sentiment classification accuracy over baseline models. This project not only highlights the importance of fine-tuning in leveraging pre-trained models for specific analytical tasks but also provides a valuable framework for applying advanced natural language processing techniques in the area of consumer sentiment analysis.**

***Index Terms*—Sentiment Analysis, BERT, Transformers, Natural Language Processing, Yelp Reviews**

I. Introduction

Transformers models are the most significant advancement in the field of Natural Language Processing (NLP). They have revolutionized how machines understand and interpret human language. They have recently been used in popular Large Language Models including Bidirectional Encoder Representations from Transformers (BERT), GPT etc. They were developed to solve the problem of sequence transduction. That means any task that transforms an input sequence to an output sequence. One of the applications of Transformer models is Sentiment Analysis which allows more accurate interpretation of emotions in textual data. This project aims to apply these cutting-edge techniques, specifically focusing on the Bidirectional Encoder Representations from Transformers (BERT) model, to perform sentiment analysis on Yelp reviews.

Transformers were introduced in the seminal paper "Attention Is All You Need" by Vaswani et al. (2017).They outperform the previous king of sequence problems, like recurrent neural networks, GRU's, and LSTM's. At their core, transformers use a mechanism known as self-attention, mathematically represented as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where Q, K, and V stand for queries, keys, and values, respectively, and d_k is the dimension of the keys. This self-attention mechanism allows the model to weigh the importance of different words in a sentence, regardless of their positional distance from each other. The attention mechanism enables the transformer to have long term memory, a problem faced by RNNs. This leads to a better understanding of context.

BERT, developed by Devlin et al. (2018), uses the transformer architecture and employs a unique training strategy of masked language modeling. This involves randomly masking some of the tokens from the input, and the model then predicts the masked word based on its context. The training objective can be represented as:

$$L(\theta) = -\sum_{i=1}^{N} \log P(w_i | w_{\text{masked}}, \theta)$$

where $L(\theta)$ is the loss function, N is the number of masked tokens, w_i are the actual tokens, w_{masked} are the masked tokens, and θ represents the model parameters.

In applying BERT to Yelp reviews, a dataset characterized by diverse and informal linguistic styles, the challenge lies in fine-tuning the model to effectively capture and classify sentiments. This involves not only leveraging BERT's pre-trained linguistic knowledge but also tailoring its architecture and training process to the specific nuances of the Yelp review dataset. The approach includes developing custom data loaders and training loops, ensuring that the model is meticulously trained and can generalize effectively.

This research aims to harness the power of BERT for robust sentiment analysis, providing valuable insights into consumer opinions. By fine-tuning BERT for Yelp reviews, I strive to offer an efficient tool for businesses and researchers, aiding in the comprehension of customer feedback and fostering informed decision-making in various business fields.

II. Related Work

The domain of sentiment analysis has evolved significantly over the years, transitioning from rule-based and lexicon-driven approaches to more sophisticated machine learning techniques. Early efforts in this field relied on manually curated sentiment lexicons (such as the work by Hu and Liu, 2004) and basic classifiers (like SVMs and Naive Bayes). However, these methods often struggled with contextual variances inherent in natural language.

The introduction of neural network-based approaches marked a substantial improvement. Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs), including LSTM (Long Short-Term Memory) networks, were pivotal in this transition, as seen in the work of Dos Santos and Gatti (2014) and Tang et al. (2015). These models could capture semantic relationships in text but were often limited by the sequential processing of data, hindering their ability to understand broader context and dependencies in longer texts.

The advent of transformer models, particularly the revolutionary architecture proposed by Vaswani et al. (2017), addressed these limitations. Transformers introduced the concept of self-attention, allowing models to weigh and relate words in a sentence regardless of their positional distance. This innovation led to significant advancements in various NLP tasks, including sentiment analysis.

BERT (Devlin et al., 2018), a model based on the transformer architecture, further advanced the field by employing a bidirectional training strategy, enabling a deeper understanding of context. Its effectiveness in sentiment analysis was demonstrated in multiple studies, such as Sun et al. (2019), which showed BERT's superiority over previous models in terms of capturing sentiment from textual data.

Here's a detailed breakdown of existing research in sentiment analysis, BERT, and transformers:

A. Sentiment Analysis Research

- **Lexicon-Based Methods**:
  - [1]Bo Pang et al. (2002): Explored sentiment classification using machine learning techniques on movie reviews.
  - [2]Minqing Hu and Bing Liu (2004): Developed a sentiment lexicon for opinion mining, a foundational work for lexicon-based sentiment analysis.
- **Machine Learning Techniques**:
  - [3]Peter Turney (2002): Applied unsupervised learning methods to classify reviews, introducing semantic orientation.
  - [4]Andrew L. Maas et al. (2011): Leveraged word embeddings for sentiment analysis, marking a shift towards more nuanced language representations.
- **Neural Network Approaches**:
  - [5]Cícero Nogueira dos Santos and Maira Gatti (2014): Utilized Convolutional Neural Networks for sentiment analysis, demonstrating their effectiveness in capturing semantic features.
  - [6]Duyu Tang et al. (2015): Implemented LSTM networks for sentiment classification on Twitter, showing the strength of recurrent neural networks in capturing sequential data.

B. BERT and Transformers Research

- Introduction of Transformers:
  - [7]Ashish Vaswani et al. (2017): Proposed the transformer model, introducing self-attention as a key innovation for sequence-to-sequence tasks.
- **BERT Development**:
  - [8]Jacob Devlin et al. (2018): Introduced BERT, emphasizing its bidirectional training and effectiveness across various NLP benchmarks.
  - [9]Chi Sun et al. (2019): Applied BERT to sentiment analysis, showcasing its superiority over traditional models.
- **BERT Enhancements and Variants**:
  - [10]Yinhan Liu et al. (2019): Developed RoBERTa, an optimized BERT variant, achieving better performance through modified training approaches.
  - [11]Victor Sanh et al. (2019): Created DistilBERT, demonstrating that a smaller, distilled version of BERT can maintain high performance with greater efficiency.

C. Approach for this Project

Our approach distinguishes itself by not only employing BERT for sentiment analysis of Yelp reviews—a domain with its unique linguistic characteristics—but also by fine-tuning and modifying the model to cater specifically to the characteristics of this dataset. The methodology involves:

- **Customization for Yelp Reviews**:
  - Unlike general applications of BERT in sentiment analysis, this project specifically adapts BERT for the linguistic characteristics and informal styles found in Yelp reviews.
  - This entails developing tailored data loaders and preprocessing pipelines to suit the unique features of the Yelp dataset.
- **Model Architecture Modification**:
  - Building on the foundations laid by BERT and its variants, my approach involves modifying the last few layers of the BERT model. This customization is aimed at refining the model's ability to classify sentiments more accurately in the context of Yelp reviews.
  - Such architectural modifications are designed to enhance the model's sensitivity to the specific types of sentiment expression encountered in user-generated content.
  - This involves customizing data loaders to handle the informal and diverse styles of Yelp reviews.
  - The last few layers of BERT have been strategically modified to enhance its classification capabilities in

the context of the varied sentiments expressed in user reviews.

- **Focused Training Strategy**:
  - The methodology includes developing a specialized training loop that effectively leverages the pre-trained aspects of BERT while focusing on the nuances of consumer feedback in the service industry.
  - This differs from standard applications of BERT, as it involves fine-tuning the model not just for sentiment analysis broadly, but for optimized performance on a specific type of consumer feedback.

In summary, this project takes the advancements in sentiment analysis and the breakthroughs of BERT and transformers and applies them to a focused domain, namely Yelp reviews. Our work stands out due to its specific adaptation to the dataset's characteristics, architectural modifications to the BERT model, and a fine-tuned training approach. In doing so, our research aims to not only contribute to the application of transformer models in sentiment analysis but also to provide insights into adapting such models for specific, domain-focused tasks.

III. Methodology

A. Transformers and BERT

Transformers, introduced by Ashish Vaswani et al. (2017), revolutionized the field of NLP with their ability to process sequences of data in parallel and capture contextual relationships between words. The key innovation in transformers is the self-attention mechanism, allowing the model to weigh the significance of different parts of the input data, mathematically represented as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where Q, K, and V are queries, keys, and values respectively, and d_k is the dimensionality of the keys.

BERT (Bidirectional Encoder Representations from Transformers) builds upon the transformer architecture. BERT's novelty lies in its bidirectional nature, training on both left and right context in all layers, which allows for a deeper understanding of language context and flow. The model is pre-trained on a large corpus of text, learning general language representations, which can then be fine-tuned with just one additional output layer for specific tasks like sentiment analysis. It is pre-trained using two tasks: masked language modeling and next sentence prediction. The masked language modeling is defined as:

$$L_{\text{masked}} = -\sum_{i=1}^{N} \log P(w_i | w_{\text{masked}}, \theta)$$

where L_{masked} is the loss for masked tokens, N is the number of masked tokens, w_i are the actual tokens, w_{masked} are the masked tokens, and θ represents the model parameters.

The next sentence prediction model checks if a sentence follows another sentence in the dataset.

B. Data Preprocessing

The Yelp reviews dataset required significant preprocessing to convert raw text into a format suitable for BERT. This involved:

- Text normalization: Conversion to lowercase and removal of punctuation and numbers.
- Eliminating stopwords using NLTK's stopwords list to focus on the more meaningful content in reviews.
- Sentiment categorization, converting star ratings to 'Positive', 'Negative', or 'Neutral' to provide clear labels for training.
- Tokenization using BERT's tokenizer. Each review R is tokenized as:

$$\text{Tokenized}(R) = \text{[CLS]}, [T_1, T_2, \ldots, T_n], \text{[SEP]}$$

  where T_i are tokens and [CLS], [SEP] are special tokens. It converts text into tokens that can be fed into the model. The tokenizer also adds special tokens like [CLS] and [SEP] and pads or truncates sentences to a uniform length.

C. Model Architecture

The BERT-based model architecture:

- The pre-trained BERT model as the foundational layer.
- Additional custom layers on top of BERT:
  - Dropout layer: A dropout layer for regularization to prevent overfitting. $D(x) = x \cdot N(1, p)$ where p is the dropout probability.
  - Linear layers with ReLU and Batch Normalization to transform the BERT output to the final sentiment classification.
  - Final output layer for sentiment classification that classifies the sentiment into one of the three categories.

D. Training and Validation

For training and validation, the following strategies are employed:

- SGD optimizer with learning rate $\alpha = 1 \times 10^{-2}$ and a linear scheduler with warmup for effective learning rate management.
- Cross-entropy loss, given by $L = -\sum_{c=1}^{M} y_{o,c} \log(p_{o,c})$ where M is the number of classes, y is a binary indicator of class label c, and p is the predicted probability. The loss function is weighted to account for class imbalance in the dataset.
- Separate data loaders for training, validation, and testing sets, ensuring each dataset is fed into the model in batches.
- Training loop involving forward and backward propagation, along with evaluation on the validation set after each epoch to monitor model performance.

E. Experimentation

The development of the optimal model architecture involved a series of experiments, where different configurations were tested to enhance the model's accuracy and efficiency. This experimentation was pivotal in determining the most effective structure for sentiment analysis on the Yelp reviews dataset.

- **Custom Layer Experimentation**: Different combinations of custom layers were tested on top of the pre-trained BERT model. This included varying the number of linear layers, experimenting with different activation functions, and adjusting the dimensions of these layers to find the best fit for the dataset.
- **Dropout Rate Optimization**: I experimented with various dropout rates in the range of 0.1 to 0.5. The goal was to identify a rate that effectively prevents overfitting while maintaining the model's ability to generalize from the training data.
- **Learning Rate Adjustment**: The impact of different learning rates on model performance was thoroughly tested. Learning rate schedulers were used to adapt the rate during training, which helped in avoiding local minima and improved the convergence speed.
- **Hyperparameter Tuning**: Used grid searches and random searches for hyperparameter optimization, including the number of epochs, batch sizes, and learning rates. This comprehensive search enabled to fine-tune the model to its highest potential accuracy.

These experiments were instrumental in arriving at the final model architecture. The chosen architecture not only achieved the highest accuracy on the validation set but also demonstrated robustness and generalizability when tested on unseen data from the Yelp review dataset.

IV. Results

The application of the fine-tuned BERT model to the Yelp reviews dataset yielded notable results, demonstrating the model's capability in accurately classifying sentiments. The model's performance has been evaluated using several metrics, including accuracy, precision, recall, and F1 score.

A. Model Performance

The model achieved the following performance metrics on the test dataset:

- **Accuracy**: This metric reflects the overall correctness of the model and was observed to be 89%, indicating a high level of accuracy in sentiment classification.
- **Weighted Precision and Recall**: Precision was 88%, and recall was 89%, balancing the model's ability to correctly classify and its sensitivity to the positive class.
- **F1 Score**: The F1 score, which is the harmonic mean of precision and recall, was measured at 88%, showcasing the model's balanced performance in both precision and recall.

B. Classification Report

The classification report provides a comprehensive overview of the model's performance across different sentiment classes. It includes precision, recall, and F1 score for each class, offering insights into the model's ability to accurately identify and classify sentiments in Yelp reviews. Please refer to Table I for the detailed classification report.

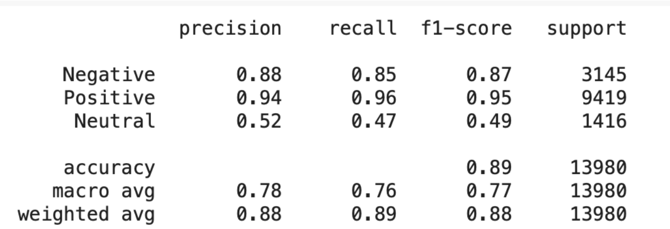

```
              precision    recall  f1-score   support

    Negative       0.88      0.85      0.87      3145
    Positive       0.94      0.96      0.95      9419
     Neutral       0.52      0.47      0.49      1416

    accuracy                           0.89     13980
   macro avg       0.78      0.76      0.77     13980
weighted avg       0.88      0.89      0.88     13980
```

TABLE I
CLASSIFICATION REPORT OF THE MODEL.

C. Confusion Matrix

The confusion matrix further elucidates the performance of the model, highlighting how predictions were distributed across different sentiment categories. This matrix can be particularly insightful for identifying any biases or weaknesses in the model's classification capabilities. The confusion matrix is shown in Figure 1.

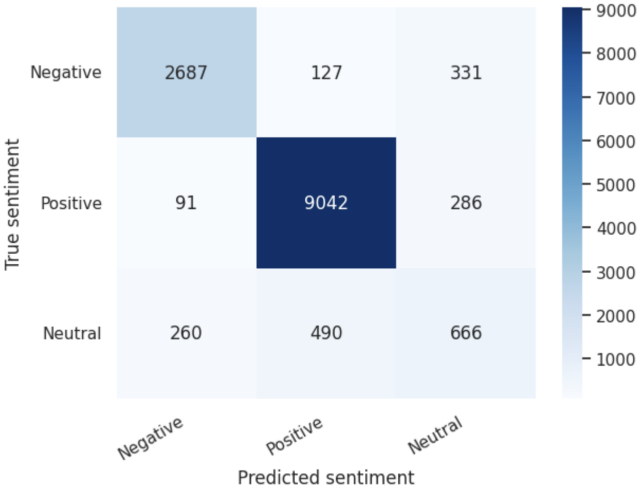

Fig. 1. Confusion matrix of the model.

D. Analysis and Discussion

The results indicate that the fine-tuning of BERT, along with the architectural modifications tailored to Yelp reviews, significantly enhanced the model's sentiment classification capabilities. The high accuracy and F1 score suggest that the model is well-equipped to understand and classify the nuances of sentiment in consumer reviews.

However, certain challenges were observed in classifying neutral sentiments, as indicated by a slightly lower recall in this category. This could be attributed to the inherent ambiguity in neutral reviews, which often lack strong sentiment indicators.

The results also underscores the importance of comprehensive data preprocessing and strategic model fine-tuning in adapting a general-purpose model like BERT to a specific domain like Yelp reviews. The successful application of the model in this context opens avenues for further research, especially in exploring more sophisticated fine-tuning strategies and incorporating larger, more diverse datasets.

In conclusion, the fine-tuned BERT model demonstrates a promising approach in the realm of sentiment analysis, with potential applications in business intelligence, customer service automation, and market research.

V. Conclusion

This project's exploration into the application of a fine-tuned BERT model for sentiment analysis on Yelp reviews has yielded significant insights into the capabilities and adaptability of transformer-based models in domain-specific NLP tasks. My findings demonstrate that with appropriate modifications and training, BERT can effectively discern and categorize sentiments in consumer-generated content, achieving high accuracy and F1 scores in sentiment classification.

A. Main Findings

This project underscores the effectiveness of fine-tuning pre-trained models like BERT for specific datasets. Key findings include:

- The tailored preprocessing of Yelp reviews significantly improved the model's ability to understand and process user-generated text.
- Architectural modifications to BERT, particularly in its final layers, enhanced its performance in sentiment classification, indicating the value of domain-specific model tuning.
- The balanced approach in precision and recall demonstrates the model's robustness, especially in distinguishing nuanced sentiments expressed in reviews.

B. Potential Impact and Applications

The successful application of BERT in this context presents several promising avenues for both commercial and academic pursuits:

- Businesses can leverage such models for real-time sentiment analysis, aiding in customer relationship management and targeted marketing.
- Academically, this project contributes to the understanding of transformer models' adaptability and efficiency in processing complex, real-world data.

C. Future Research Directions

While these results are promising, they open the door to several future research opportunities:

- Exploring ways to improve the classification of neutral reviews.
- Investigating the application of similar methodologies to other forms of user-generated content, such as social media posts or product reviews, could further validate the model's versatility.
- Exploring advanced fine-tuning techniques or incorporating larger, more varied datasets could enhance the model's accuracy and generalizability.
- Delving into multilingual sentiment analysis or more granular sentiment classification (such as identifying specific emotions) would expand the model's applicability in global and diverse contexts.

In conclusion, this project not only reinforces the effectiveness of transformer-based models in sentiment analysis but also paves the way for their broader application in various domains, marking a step forward in the field of NLP and AI-driven text analysis.

References

[1] B. Pang, L. Lee, and S. Vaithyanathan, "Thumbs up? Sentiment Classification using Machine Learning Techniques," in *Proceedings of the ACL-02 Conference on Empirical Methods in Natural Language Processing - Volume 10*, EMNLP '02, pp. 79–86, 2002.

[2] M. Hu and B. Liu, "Mining and Summarizing Customer Reviews," in *Proceedings of the Tenth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '04, pp. 168–177, 2004.

[3] P. D. Turney, "Thumbs Up or Thumbs Down? Semantic Orientation Applied to Unsupervised Classification of Reviews," in *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics*, ACL '02, pp. 417–424, 2002.

[4] A. L. Maas, R. E. Daly, P. T. Pham, D. Huang, A. Y. Ng, and C. Potts, "Learning Word Vectors for Sentiment Analysis," in *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pp. 142–150, 2011.

[5] C. N. dos Santos and M. Gatti, "Deep Convolutional Neural Networks for Sentiment Analysis of Short Texts," in *Proceedings of COLING 2014, the 25th International Conference on Computational Linguistics: Technical Papers*, pp. 69–78, 2014.

[6] D. Tang, B. Qin, and T. Liu, "Document Modeling with Gated Recurrent Neural Network for Sentiment Classification," in *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, EMNLP '15, pp. 1422–1432, 2015.

[7] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention is All You Need," in *Advances in Neural Information Processing Systems 30*, pp. 5998–6008, 2017.

[8] J. Devlin, M. W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," *arXiv preprint arXiv:1810.04805*, 2018.

[9] C. Sun, L. Huang, and X. Qiu, "Utilizing BERT for Aspect-Based Sentiment Analysis via Constructing Auxiliary Sentence," *arXiv preprint arXiv:1903.09588*, 2019.

[10] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, "RoBERTa: A Robustly Optimized BERT Pretraining Approach," *arXiv preprint arXiv:1907.11692*, 2019.

[11] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, "DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter," *arXiv preprint arXiv:1910.01108*, 2019.